Inshallah Corpus Collection and Cleaning

Phase 2 report

Scope

D2.1 contains 479,942 occurrences of the formula. D2.2 contains 1,000 rows: 300 SFC, 350 Gulf and 350 SADA. label and context_needed are left empty for Phase 3.

Sources and access

Source	Records read	Files	Recorded licence
SFC mini [1]	8,866,859	63 text files	CC BY NC 4.0
Gulf tweets [2]	108,052	1 CSV	CC BY 4.0
SADA text [3]	253,166	train valid test CSVs	CC BY NC SA 4.0
Tanzil [4]	6,236 verses	simple clean text	CC BY 3.0 with verbatim terms

Published distributions were downloaded on 9 October 2026 without new scraping. Gulf text is publisher-cleaned; SADA uses GroundTruthText without audio. Tanzil supplies the quotation reference. SFC/Gulf revisions, SADA version 23, source hashes and licences are recorded in the access log.

Extraction and cleaning procedure

We search a normalized copy of each text, in which diacritics and tatweel are removed and the alef and hamza forms are made the same. The search finds the formula written with spaces, joined or with extra letters, and records where each match starts and ends in the original text. Each match becomes its own row. original_text is the full record exactly as it was downloaded, and variant is the matched part copied as it is.

Sentence boundaries identify the host sentence. Forum and tweet prev_context comes only from the same record, never a reply thread. For SADA, context is the preceding time-ordered segment in the same recording. A missing predecessor produces an empty field. Links are protected during sentence splitting, then links and user mentions are removed from the working host and context, without changing original_text.

We removed duplicates before sampling. Exact copies were removed first. For near-copies, MinHash LSH suggested pairs, and a pair counted as a duplicate only if their words overlapped by at least 90% (Jaccard 0.90). When a record is kept, all of its formula occurrences are kept, so a post that uses the formula twice keeps both rows. LSH can miss some near-copies.

Source	Found	After duplicates	After all filters
Gulf	1,036	926	885
SADA	9,655	9,627	9,114
SFC	519,971	508,010	469,943
Total	530,662	518,563	479,942

The counts are occurrences of the formula, not posts or speakers. Removing hashtag-only posts took out no further rows after duplicate removal. The length rule keeps host sentences of 4–80 words, counted with the CAMeL tokenizer without punctuation or emoji. Quotations stay in D2.1: 208 occurrences have quote_flag=true, so 479,734 occurrences could be sampled.

Quotation and descriptive rules

To find Quran quotations, we compare normalized word windows with the Tanzil text; a match needs at least two Quranic words besides the formula. A short list of fixed prayers covers other quotations. The flag belongs to one occurrence, so a prayer somewhere else in a long sentence does not remove it.

position is first, middle, last or alone, measured from the marked occurrence. verb_cue looks at the three words before and after it, using CAMeL analyses and a list of exceptions. future_marker looks for future particles and likely verb prefixes, including the Arabic s prefix; some nouns and participles look similar, so they need checking. second_person looks for pronouns and endings that address someone, with exceptions such as speaking to God. These columns are rough cues, not labels.

Output schema and occurrence provenance

D2.1 uses the 11 prescribed columns, D2.2 adds stratum, label and context_needed, and a separate provenance file records source references, original hashes, offsets and stable occurrence keys.

Descriptive statistics from the full collection

Numbers in this section come from all of D2.1, not from the 1,000-row sample. D2.1 has 1,095 different spellings (1,222 before cleaning), because extra letters and diacritics make a new spelling. The full lists are in table_variants_total.csv and table_full_cleaned_variants.csv.

Exact surface variant	Occurrences
ان شاء الله	271,716
إن شاء الله	75,780
ان شاءالله	30,832
انشاء الله	23,897
انشالله	13,934
ان شالله	12,195
أن شاء الله	12,010
انشاءالله	4,843
إن شاءالله	4,172
ان شا الله	3,105
Other 1,085 exact variants	27,458

Field	Category counts in D2.1
position	first 55,058; middle 331,994; last 92,888; alone 2
verb_cue	me_future 81,850; me_past 21,912; other 270,699; none 105,481
future_marker	true 73,446; false 406,496
second_person	true 288,047; false 191,895
quote_flag	true 208; false 479,734

Source	Mean words	Median	Range	With previous context
Gulf	23.6	20	4–69	15 of 885 • 1.7%
SADA	30.5	25	4–80	8,934 of 9,114 • 98.0%
SFC	22.7	17	4–80	124,809 of 469,943 • 26.6%

The four-word minimum removes almost all standalone matches (only two remain); short_or_alone and answers_question therefore replace the guide’s standalone sampling stratum. Context availability also differs sharply by source; per-source descriptive tables accompany this report.

Human checking and correction evidence

We checked the 50-row metadata sheet starting from AI suggestions and decided each value ourselves; earlier corrections that matched were kept. We checked a random sample of 50 flagged quotations in the same way and kept one uncertain case. Neither check is independent expert annotation.

Field	Checked	Incorrect before repair	Rate
variant	50	0	0%
position	50	0	0%
verb_cue	50	13	26%
future_marker	50	6	12%
second_person	50	4	8%
quote_flag	50	0	0%

The 50-row check found 23 wrong values: 13 verb_cue, six future_marker and four second_person. We corrected them on the checked rows; this does not show the accuracy on unchecked rows. The correction log records 39 changes in total.

To check the quotation flag, we picked 50 of the 208 flagged occurrences at random (seed 2026). Seven are real Quran quotations, 42 are not, and one is unclear, so only 7 of 50 flags are right (14%), or 8 of 50 if the unclear one is counted. Most wrong flags are ordinary uses next to a fixed prayer. Every decision is in D2_quote50_audit.csv. We did not change the flags or D2.2, and this check does not tell us how many quotations were missed.

Independent random planning check

The random planning check has 30 SFC, 35 Gulf and 35 SADA rows. We gave each row a quick label; we started from AI suggestions and checked each one. Multiplying the counts by ten gives an estimate for 1,000 rows with this source mix, not for the whole corpus.

Function	In 100	Estimate in 1,000	Planning interpretation
COMMIT	17	170	At least 50
WISH	58	580	At least 50
DEFER	1	10	Below 50
HEDGE	13	130	At least 50
SKEPTIC	4	40	Below 50
CLOSE	1	10	Below 50
OTHER	3	30	Outside six target labels
UNCLEAR	3	30	Outside six target labels

We also checked 312 candidate rows found by targeted searches, again starting from AI suggestions. They guide sampling only; they are not a random sample or gold labels. The rates used for sampling in the next section come from the random 100 rows and D2_stratum_extra_filled.csv.

Selection of the annotation sample

Allocation keeps the source quotas and at least one row from each populated stratum. The evidence does not support 60 expected examples for both CLOSE and DEFER while keeping 50 for the other four labels (margin_feasibility.json).

The 312 candidates supply verified strata (12 CLOSE, seven DEFER and one SKEPTIC). The table below shows the largest and the rarest strata; table_strata_evidence.csv lists all of them.

Stratum	Draw	Reviews	DEFER / SKEPTIC / CLOSE	Evidence
sfc_skeptic_cue	282	14	7.1% / 7.1% / 0%	Local random + extra
gulf_other	334	14	0% / 7.1% / 0%	Local random100
sada_addressed	332	26	7.7% / 7.7% / 11.5%	Local random + extra
sfc_verified_close	7	7	0% / 0% / 100%	Candidate sheet
sada_verified_close	5	5	0% / 0% / 100%	Candidate sheet
sfc_verified_defer	2	2	100% / 0% / 0%	Candidate sheet
sada_verified_defer	5	5	100% / 0% / 0%	Candidate sheet

The three largest cells hold 94.0%, 95.4% and 94.9% of their source samples; we chose this on purpose to get more of the rare labels.

Core function	Expected in final sample	Meets minimum
COMMIT	167.26	Yes
WISH	555.50	Yes
DEFER	53.65	Yes
HEDGE	75.33	Yes
SKEPTIC	71.21	Yes
CLOSE	50.62	Yes

These are planning estimates, not real labels. CLOSE (0.62) and DEFER (3.65) are only just above 50. If either has fewer than 40 natural rows after annotation, we will use the Phase 3 fallback.

Each of the 1,000 IDs is different, appears in D2.1 and has quote_flag=false, and the annotation columns are empty. Every stratum name starts with its source, such as gulf_other and sfc_skeptic_cue. We sampled with seed 2026 without replacement and shuffled the rows at the end; the weights come with the sample. 31 rows in D2.2 come from 15 sentences that contain the formula more than once; each row is a different occurrence, identified by its position in D2.1_provenance.csv.

Datasheet limitations and intended use

This collection is intended for research on formula functions in Saudi-associated public text. It is an occurrence corpus, not a survey of Saudi speakers. Source membership does not prove the nationality or dialect of each author. Gulf material is heavily COVID-related and includes broader regional Arabic. SFC makes up 97.9% of D2.1, while D2.2 uses a different source mix on purpose.

SADA has acted dialogue and noisy transcripts with several speakers. Of 9,114 retained occurrences, 5,612 have a multi-speaker dialect tag, 1,669 Najdi, 673 Hijazi, 564 Khaliji and 471 Unknown. Named non-Saudi tags also remain, including 53 Egyptian and 20 Levantine occurrences. No strict speaker-level Saudi-only claim is made. The previous segment may come from another speaker, and missing punctuation or broken words can hide what the speaker meant.

The 4–80 word rule removes most short acknowledgements and long stories. The morphology and quotation rules are rough. The quotation check shows that many flags are wrong, and we do not know how many quotations were missed. original_text keeps the full record and may contain personal information.

Reproducibility and verification

The executed notebook ran without errors. The README explains where to put the sources and manual sheets, the order of the cells and the output files.

Final checks confirmed the schema, the source quotas, unique IDs, empty annotation fields, no quotations in D2.2, every D2.2 row present in D2.1, and the 4–80 word bounds. All 479,942 original records and formula offsets were checked against the downloaded files.

Repository: https://github.com/Nuha-Alodan/inshallah-phase1. D2.1 (396 MB) is too large for GitHub and is shared here: https://drive.google.com/file/d/1wJbdsxksV2r_piTCvYuox2JU4MDstBob/view?usp=drive_link

References and companion evidence

[1] SFC mini dataset card and licence. https://huggingface.co/datasets/faisalq/SFC-mini

[2] Gulf Arabic Tweets 2018–2020 dataset card and licence. https://huggingface.co/datasets/AhmedSSabir/Gulf-Arabic-Tweets-2018-2020

[3] SADA Kaggle text distribution version 23 and dataset card. https://www.kaggle.com/datasets/sdaiancai/sada2022/versions/23 ; https://huggingface.co/datasets/MohamedRashad/SADA22

[4] Tanzil text licence and download. https://tanzil.net/docs/text_license ; https://tanzil.net/download/

Run evidence: tables and run_manifest.json in reports/phase2; reviewed sheets in manual; sampling weights in data; D2.1 provenance with D2.1. Source pages checked 9 October 2026.